

EECS 370 Final Cheat Sheet

ISA

- RISC : Reduced instruction set computing
- CISC : Complex instruction set computing

ARM, MIPS & LC2K are both RISC. All RISC instructions have the same length (eg: 32-bits for LC2K). Intel x86 has CISC interface, with RISC core.

MIPS

```
add $s1, $s2, $s3    $s1 = $s2 + $s3
beq $s1, $s2, 25      if($s1==$s2) pc=pc+4+100
lw $s1, 20($s2)       $s1 = Memory[$s2 + 20]
lui $s1, 20           $s1=20*216
Sign extend if signed, and loading half-word or byte. If 0x100
= 00, 0x101 = FF, lbu $s1, 100($s0), then $s1 = 0x000000FF.
```

LC2K

```
add regA regB destReg  destReg=regA+regB
nor regA regB destReg  destReg=nor(regA, regB)
lw regA regB field     regB=*(regA+offset)
sw regA regB offset    *(regA+offset)=regB
beq regA regB offset   if(regA==regB)pc+=(offset+1)
jalr regA regB         regB=pc+1;pc=regA;
halt                  stop program
noop                  nothing
```

Endianess

- Big-endian: most significant bits stored first
- Little-endian: least significant bits stored first

Two's complement

Most significant bit is negative, all else is positive.

$$10000001_{\text{two}} = -1 \cdot 2^7 + 1 = -127$$
$$00000001_{\text{two}} = -1 \cdot 2^0 + 1 = 1$$

To switch signs, invert and add 1. Also assume unlimited sign extension of most significant bit (eg: $1001_{\text{two}} = \dots 111001_{\text{two}}$)

C to Assembly

Memory layout

Golden rule

Address of variable is aligned based on size of the variable.

- char is byte aligned (any addr)
- short is half-word aligned (least significant bit == 0)
- int is word aligned (two least significant bit == 0)
- pointer is size of computer's architecture

Structs

- starting addr % size(largest basic type) == 0
- size % size(largest basic type) == 0
- allows array of structs

```
struct record {
    float a;          //0-3
    char b;            //4-4, 5-7 padding
    struct {
        char ** c;    //8-11
        short d[3];   //12-17, 18-19 padding
    } s;              //8-19, 12 % 4 == 0
                      //19-23 padding
    double e;         //24-31
    short * f;        //32-35, 37-39 padding
};                   //0-39, 40 % 8 == 0
```

Function calls

Pushed & popped on call-frame stack in following order:

1. incoming params
2. \$fp
3. return address
4. callee saved registers (don't save if not needed)
5. local vars
6. spilled registers (when not enough registers)
7. caller saved registers (don't save dead variables)
8. outgoing parameters
9. \$sp

Liveness example

```
int bar(void) {
    int d = 9;          // d: dead,
    printf("1");        // overwrite on line 4
    d = 5;              // d: alive
    int e = d;          // e: alive
    printf("2");
    return d + e;       // d,e: dead
}
```

Initialization example

```
int g = 0;              // g: static
void bar(int p, int q) { // bar: text
    return p + q;        // p, q: stack
}
int foo() {              // foo: text
    int a = 0;           // a: stack
    static int b = 1;    // b: static
    void (*func_ptr)(int, int) = &bar;
    // func_ptr: stack
    // *func_ptr: text (points to function)
```

```
int *b_ptr = &b;
// b_ptr: stack
// *b_ptr: static (points to global)
static int *m_ptr = malloc(sizeof(int));
// m_ptr: stack
// *m_ptr: heap (dynamic allocated)
return a;
}
```

Translation

header	size of other parts
text	machine code
data	globals & statics
symbol table	symbols & values
relocation table	references to variable addresses
debug info (optional)	map to source code

data does not contain uninitialized data, but keeps track of how much is needed.

symbol table contains:

- stuff other files need (globals & funcs defined in this file)
- stuff you need (referenced globals & funcs in this file defined elsewhere)
- static variables

relocation table contains:

- address the moved instruction
- type of instruction (lw, sw, jal, etc)
- referenced symbol

```
1 int brown_dog = 656;
2 short black_dog = 343;
3 extern void play_fetch(void);
4 extern int my_dog_age;
5
6 int main() {
7     if(my_dog_age == 0)
8         static int squirrel = 0;
9     my_dog_age = black_dog;
10    int i = 0;
11    while(i < 3) {
12        play_fetch();
13        I++;
14    }
15
16    int dog_treat = 3;
17    int dog_treat_bag = dog_treat * 55;
18    my_dog_age = brown_dog;
19    return 0;
20 }
```

Symbol Table		Relocation Table		
squirrel	Data	8	store	squirrel
brown_down	Data	9	store	my_dog_age
black_dog	Data	9	load	black_dog
play_fetch	Undefined	12	branch	play_fetch
my_dog_age	Undefined	18	store	my_dog_age
main	Text	18	load	brown_dog

Loader

- Creates large enough address space for program to hold text, data, stack
- Copies instructions & data from executable file memory
- Initializes registers (PC and SP most important)

Overview

1. Compiler: $*.c \rightarrow *.s$ (Assembly)
2. Assembler: $*.s \rightarrow *.o$ (1st pass: Machine code, 2nd pass: Label resolution)
3. Linker: $*.o \rightarrow *.o$, resolves absolute addresses
4. Loader: executes result

Floating math

$$\begin{aligned}
 596.75 &= 59675/100 = 2387/2^2 \\
 &= 100101010011_{\text{two}} \cdot 2^{-2} \\
 &= 1.00101010011_{\text{two}} \cdot 2^9 \\
 &= (-1)^0 \cdot 2^{(136-127)} \cdot (1 + 0.00101010011) \\
 &= (-1)^{\text{sign bit}} \cdot 2^{\text{exponent}-127} \cdot (1 + \text{mantissa})
 \end{aligned}$$

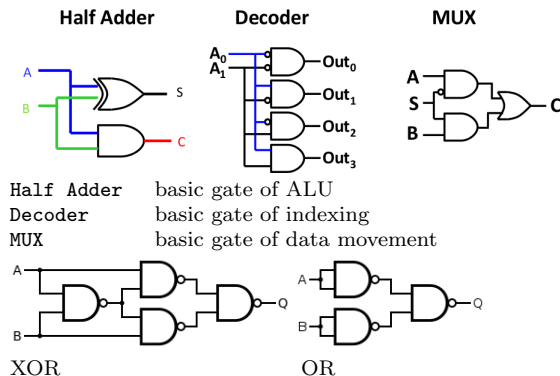
Addition is also the same, raise lower exponent to higher exponent. (Lose least significant bits)

Processor design

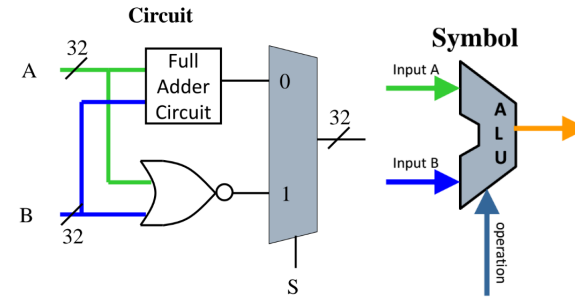
ROM bits = $2^{\# \text{ of input bits}} \cdot \# \text{ of output bits}$

Combinational logic

Combinational circuits



LC2K-ALU



For LC2K, operation may either be add or nand.


Sequential logic

Transparent D-latch

Q is set to D only if gating signal (G) is 1, else the original value is maintained. In finite-state machines, think of D as next state and Q as current state.

Edge triggered D flip-flop

positive edge: edge moving from 0 to 1
negative edge: edge moving from 1 to 0
 Made using two d-latches, with direct/inverted clock signal going to two latches's gating signal.
inverter on first latch: positive edge activated
inverter on second latch: negative edge activated



Tri-state logic

Three possible states: one, zero, not connected

Finite state machine

$T(s, i)$: current state & input to output state
 $P(s)/P(s, i)$: current state to output
 Types of FSM:
 Moore ($P(s)$): Output dependent on curr state
 Mealy ($P(s, i)$): Output dependent on curr state & input

Single cycle processor design

1. Fetch instructions
2. Decode instructions (Read data from ROM)
3. ROM output controls data movement (inc. PC, read registers, ALU control)

Each instruction is one cycle, clock drives movement. Clock period is the time it takes to execute the slowest instruction.

Clock speed example

Instr	Instr Mem Access	Read Reg	ALU	Data Mem Access	Write Reg
add	✓	✓	✓		✓
nor	✓	✓	✓		✓
lw	✓	✓	✓	✓	✓
sw	✓	✓	✓	✓	
beq	✓	✓	✓		
jalr	✓	✓			✓
noop	✓				
halt	✓				

Read memory	4 ns
Write memory	3 ns
Read register file	5 ns
Write register file	4.5 ns
ALU	4.5 ns
Others	0 ns

Then slowest instruction would be lw, since

Read instruction	4 ns
Read regA & regB	5 ns
ALU regA + offset	4.5 ns
Read resulting memory	4 ns
Write regB	4.5 ns
Total	22 ns


Control building blocks

MUX: Select either one of input
Decoder: Map instruction to ROM address

Compute building blocks

ALU: Basic arithmetic operations (ADD or NAND)
Sign extension: Repeats most significant bit to output size
 $OUT(31:0) = SE(IN(15:0))$
 $OUT(31:16) = IN(15)$
 $OUT(15:0) = IN(15:0)$

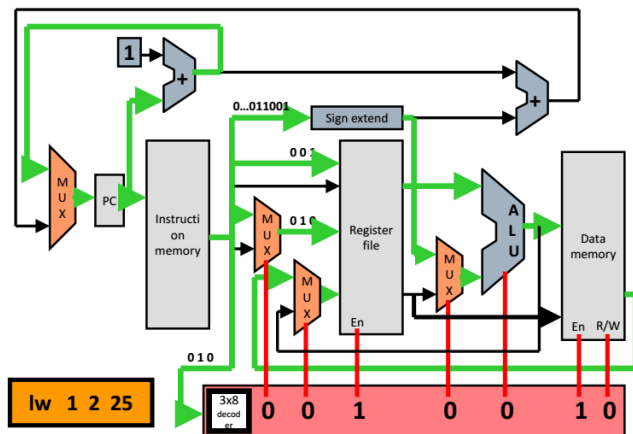
State building blocks



R1, R2, W are 3 bits each, specify OUT1 & OUT2.

LC2K datapath example

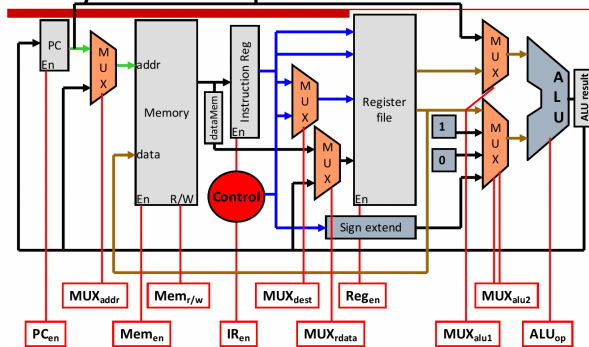
Executing a **LW** Instruction on LC2Kx Datapath



Multi-cycle CPU

Multi-cycle datapath example

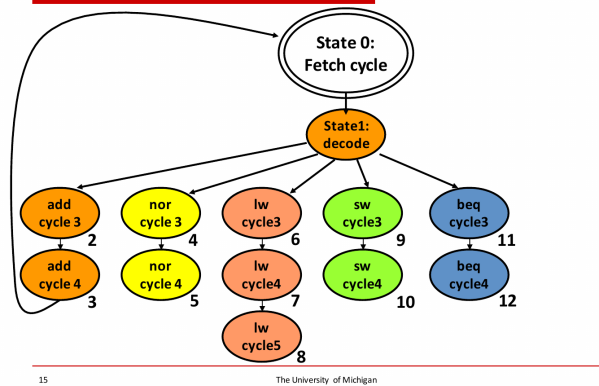
Multicycle LC2Kx Datapath



The University of Michigan

Multi-cycle FSM example

State machine for multi-cycle control signals (transition functions)



FSM example

0. Fetch Cycle

(a) $\text{mem}[\text{PC}] \rightarrow \text{instruction register}$

(b) Calculate $\text{PC} + 1$

1. Decode

2. ALU

(a) $\text{regA} + \text{offset}$ for lw

(b) for beq, do equality check also

3. Read memory (Can store in a temporary data mem)

4. Write back (Write actual value to register)

Performance metrics

• Response time (Execution time)

- Execution time = total instructions \times CPI \times clock period
- CPI = avg number of clock cycles per instruction for an application, For multicycle we need
 - Cycles necessary for each type of instruction
 - Mix of instructions executed in the application
 - Calculate new CPI from baseline CPI

• Throughput (work/time)

Pipelining... *for dummies!*

Hazards

- Data:** since register reads occur in stage 2 and register writes occur in stage 5 it is possible to read the wrong value if it is about to be written.
- Control:** A branch instruction may change the PC, but not until stage 4. What do we fetch before that?
- Exception:** How do you handle exceptions in a pipelined processor with 5 instructions in flight?

Data hazards

- Avoid:** Insert noops
 - not portable
 - large programs
 - CPI is 1, but execution is slower due to noops

```

add 1 2 3    write 3 in cycle 5
noop
noop
nand 3 4 5   read 3 in cycle 5
    
```

• Detect & Stall: CPI increases

- pass noop to execute
- stall at decode, if RaW dependency

• Detect & Forward: May still need stalling.

```

lw 3 6 10    reg 6 at MEM stage
sw 6 2 10    stall 1 cycle
    
```

Control hazards

- Avoid**
 - not portable
 - large programs
 - CPI is 1, but execution is slower due to noops

• Detect & Stall: CPI increases

- if opcode is branch, pass noop to decode

• Speculate & Squash

- assume not equal
- squash at write-back, draw table
- send noop to memory, execute, decode
- branch direction prediction using FSM

Cache Magic

- Temporal locality:** if you access a memory location (e.g., 1000) you will be more likely to re-access that location than you will be to reference some other random location
- Spacial locality:** if we reference a memory location (e.g., 1000), we are more likely to reference a location near it (e.g. 1001) than some random location

Cache avg latency

- Avg access latency = cache latency \times hit rate + memory latency \times miss rate
- Cache: 1 cycles (90% hit)
- Memory: 36 cycles (99% hit)
- Disk: 400 cycles (100% hit)

Latency: $1 + 0.1 \cdot (36 + 0.01 \cdot 400) = 5$

Stores?

- In cache? Send it to cache.
 - write-through: also send it to memory
 - write-back: mark it dirty, write on evict
- Not in cache? Write to memory
 - allocate-on-write: bring the memory to cache first?
 - no-allocate-on-write: don't do that

Mapping

- Fully-associative: Any memory location can be copied to any cache line.
- Direct-mapped: Memory block can go into specified cache block. Line index bit size is $\log_2(\# \text{ of blocks})$
- Set-associative: Like directed-mapped but with fewer sets. Couple of blocks make a set. Set index bit size is $\log_2(\# \text{-way})$. $\# \text{ sets} = \# \text{-lines} / \# \text{-ways}$

Misses

- Compulsory: First reference to a block
- Capacity: Run out of space
- Conflict: Block is evicted

How? Three step process!

1. Simulate cache with ∞ size \rightarrow compulsory misses
2. Fully-associative cache with intended size \rightarrow any new misses capacity misses
3. Actual cache \rightarrow any new misses conflict misses

Virtual Memory

Uses:

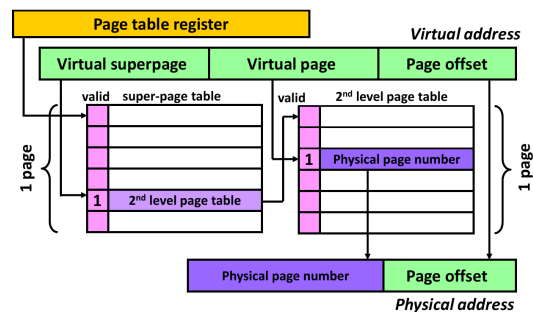
- Multiple program can share memory without:
 - transparency: know other programs exist
 - protection: program only access its' memory
- Page fault: V pages misses, handled as exceptions
- Page table: virtual page \rightarrow physical page, get physical address by physical page with offset

Page table

1GB = 2^{30} bytes of physical memory & 4KB = 2^{12} page size, then the physical page number is $30-12 = 18$ bits, plus another valid bit + other useful stuff (read only, dirty, etc.). Approx 3 bytes. How can we organize it?

- Continuous 3MB region of physical memory
- Use a multi-level page table! (Build a hierarchical page table, keep in physical memory only the translations used)
 - Super page table in physical memory
 - Second (and maybe third) level page tables only if needed
 - Size is proportional to the amount of memory used
 - Example: Assuming size is 4KB
 - * Min memory used: 4KB (Super page table)
 - * Max memory used: 4KB + $1024 \cdot 4\text{KB}$

Hierarchical page table



The University of Michigan

26

How can we replace it?

- LRU: Reference bit on page, OS clears occasionally. Evict any unreferenced page.

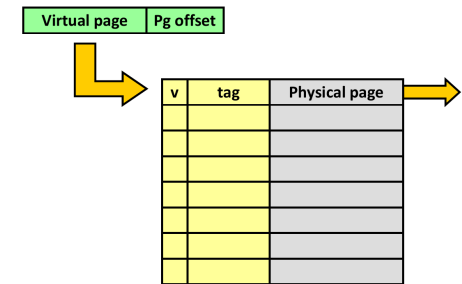
VM cache

- Virtual addr: Faster access? More complex?
- Physical addr: Delayed access? More complex?

Translation look-aside buffer (TLB)

- Avoid main memory in virtual \rightarrow physical addr translation (so it's fast)

TLB



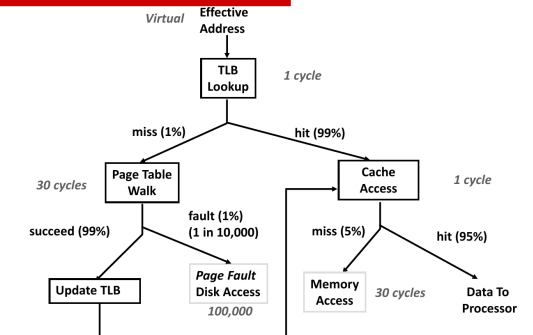
EECS 370: Introduction to Computer Organization

The University of Michigan

28

Class problem -

Physically Addressed



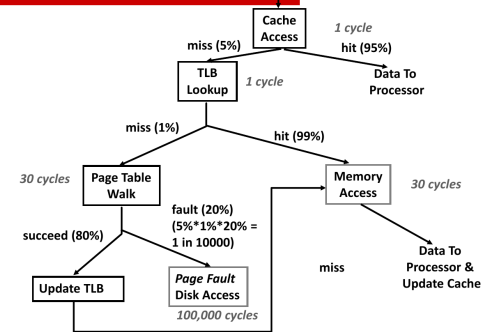
TLB + hit case + miss case

$$1 + 0.99 * (1 + .05 * 30) + .01 * [30 + 0.99 * (1 + .05 * 30) + .01 * 100,000]$$

57

Class problem -

Virtually Addressed



cache access + miss case

$$1 + .05 * [1 + 0.99 * 30 + .01 (30 + 0.8 * 30 + .2 * 100,000)]$$

59